

Data center resource management for in-network processing

Marco Micera

Politecnico di Torino, Technische Universität Darmstadt

*“Modern Internet services, such as search, social networking, and e-commerce, **critically depend on high-performance key-value stores**. Rendering even a single web page often requires hundreds or even thousands of storage accesses.”*

NetChain [2] authors

*“As the number of compute elements grows, and the need to expose and utilize higher levels of parallelism grows, **it is essential to [...] focus on developing architectures that lend themselves better to providing extreme-scale simulation capabilities.**”*

SHArP [1] authors

In-Network Processing (INP)

- INP refers to the technique of **offloading parts of the computation to network devices** (e.g., programmable switches, network accelerators, middleboxes, etc.), hence reducing the load on servers
- Advantages:
 1. Serve network requests on the fly with low latency
 2. Reduce data center traffic and mitigate network congestion
 3. Save energy by running servers in a low-power mode
- Few solutions out there already: Daiet [4], SHArP [1], NetChain [2], IncBricks [3]

Problem statement

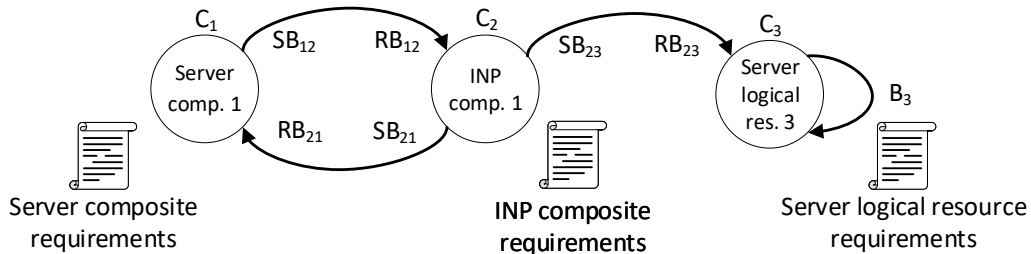
For the time being, it seems that there is still no Resource Manager (RM) that takes into account the presence of a network having a data plane that supports (partially or completely) INP

Goals

1. Model and evaluate an API through which applications can ask for INP resources
2. Discuss the importance of a scheduler which can reject INP requests and propose their server-only equivalent when needed (e.g., high switch utilization)

Design

The Extended-Tenant Application Graph (eTAG)



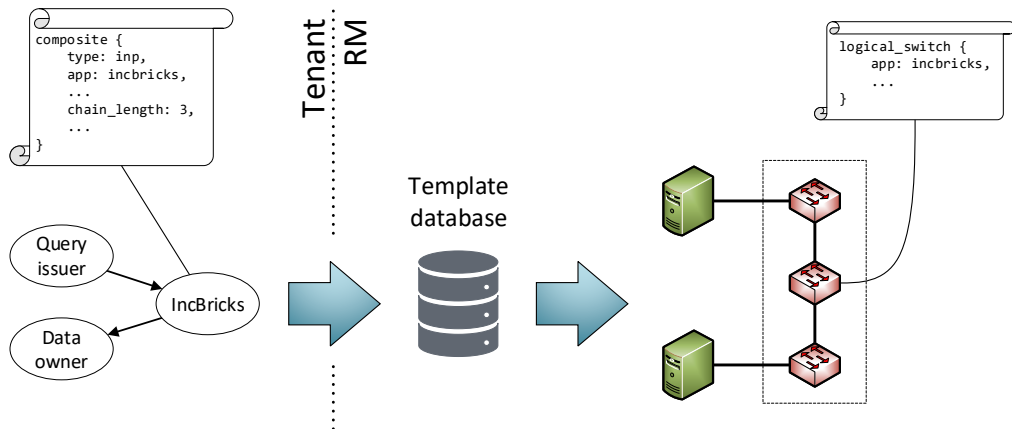
- A composite is a template that describes high-level logical components
 - It can be of two types:
 - Server (e.g., “web server”, “database”, ...)
 - INP (e.g., “IncBricks caching system”, “NetChain locking system”, ...)
 - It can be made out of
 - Other Composites
 - Logical resources

Generic groups

- In-network **storage**
 - Switches must
 - dedicate part of their local memory to store a distributed map
 - form a chain
 - IncBricks [3], NetChain [2]
- In-network **data aggregation**
 - Switches must
 - form a tree whose root is connected to data consumers and whose leaves are connected to data producers
 - dedicate part of their local memory to store a key-value map
 - be able to perform basic operations on data, such as writing and hashing
 - wait for all its children to send aggregated data
 - Dalet [4], SHArP [1]

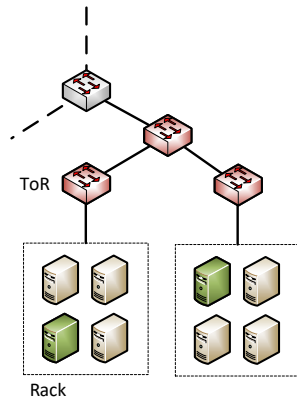
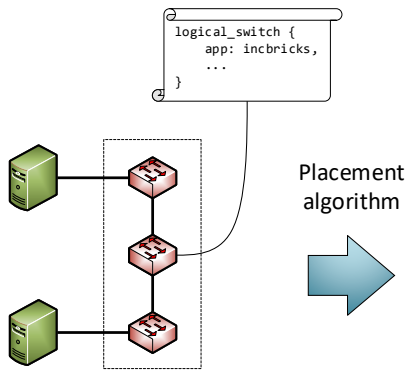
Mapping composites to logical resources

- The *template database* maps composites (or generic groups) to their equivalent made out of just logical resources

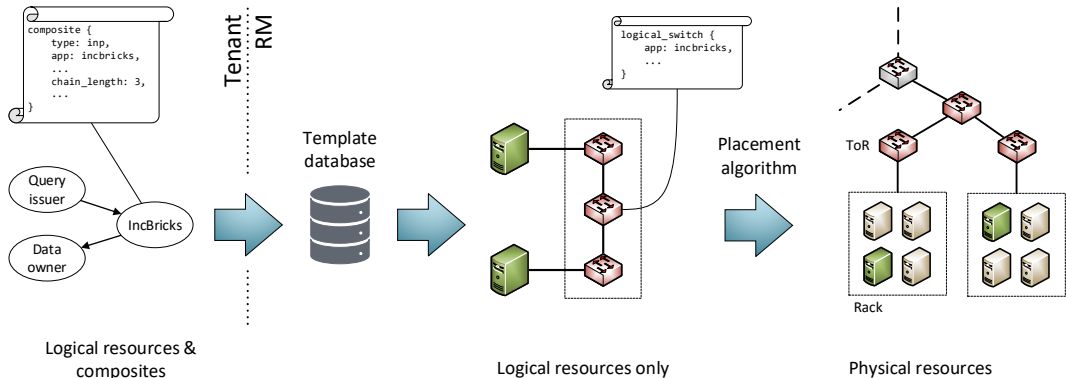


Placing logical resources

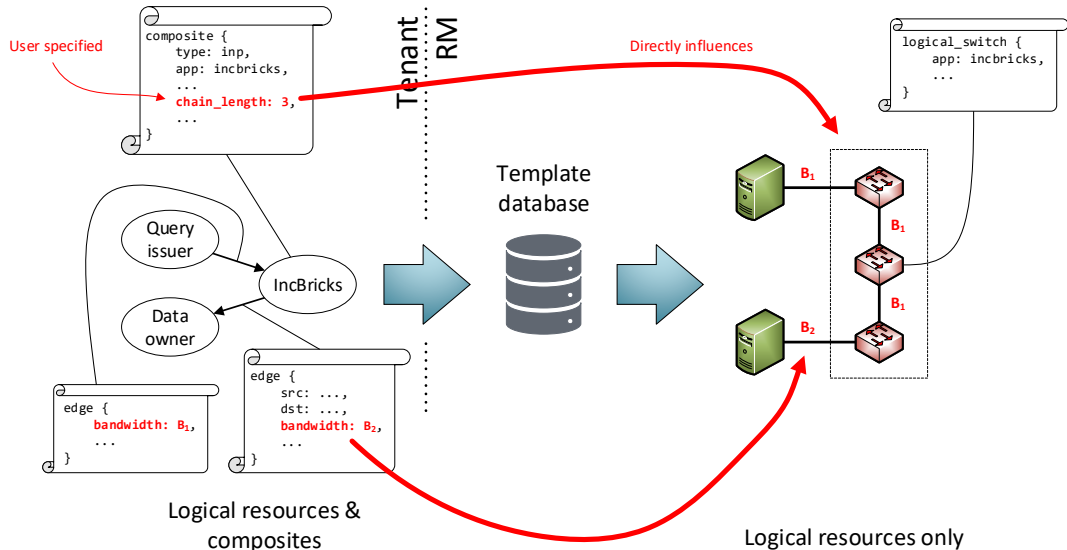
- The placement algorithm places logical resources onto physical ones
- Three types of physical resources:
 - Server
 - Switch
 - Link



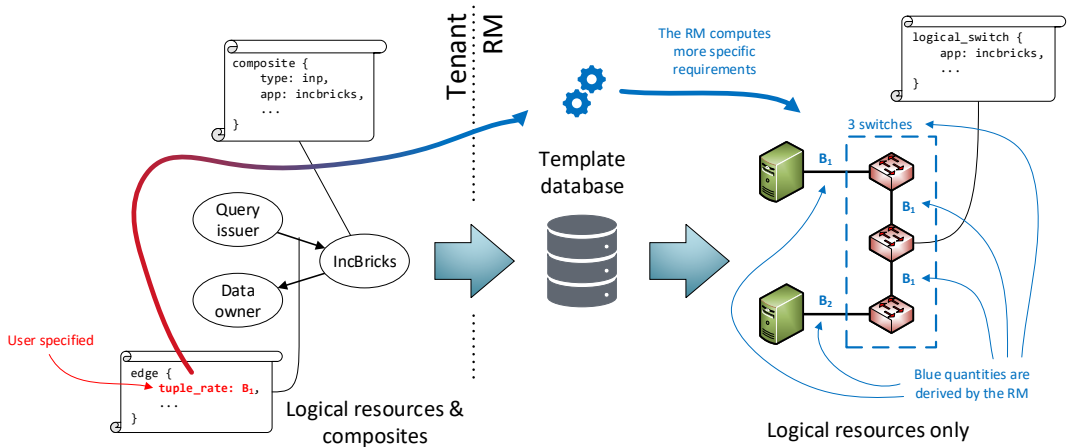
The whole picture



1st approach: passive template mapping



2nd approach: active template mapping



Evaluation

Simulation 1/2

- Simulator built from the ground up
 - Inspired by Omega's [5] *lightweight simulator*²
 - Supports multiple resource dimensions, switch resources, and composites
 - Switches have *properties* (e.g., supported INP solutions)
- Simulated data center physical architecture: fat-tree with 4 pods
- 3 days-long randomly-generated workload
 - Job properties (e.g., requirements, requests' interarrival time, etc.) are sampled from exponential distributions
- Simple greedy scheduler

²Available at github.com/google/cluster-scheduler-simulator

- The template database contains two entries for the previously-mentioned generic groups
 - In-network storage (switch chain)
 - In-network data aggregation (switch tree)
- Server Tasks Cutback (STC): the reduction of server tasks once an INP solution is introduced

$$STC = \frac{\#server\ tasks\ without\ INP}{\#server\ tasks\ with\ INP} \quad (1)$$

- Sweep: percentage of requests including INP composites

Results 1/3

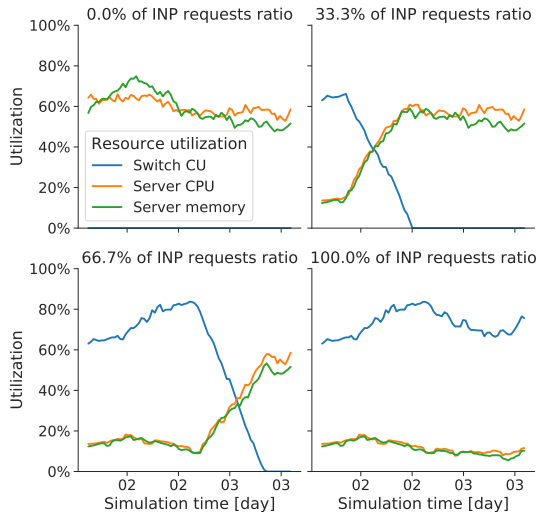


Figure 1: physical resource utilization for different amounts of INP requests

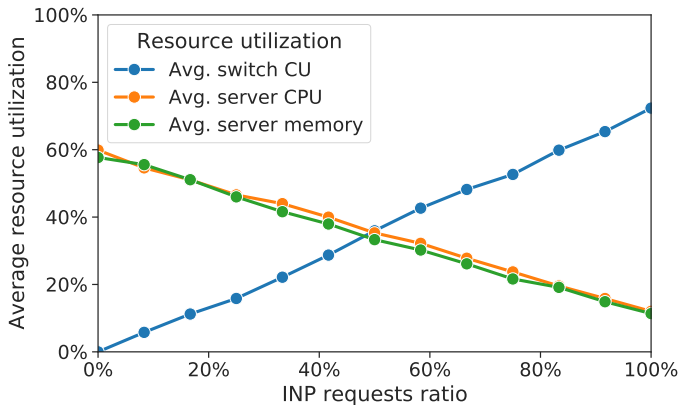


Figure 2: Average resource utilization as a function of the INP requests ratio

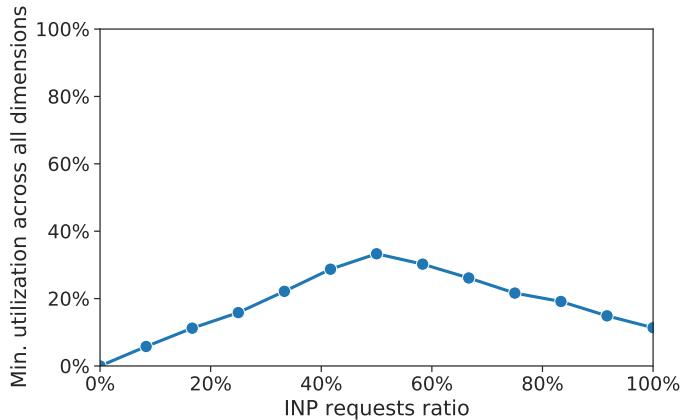


Figure 3: Minimum resource utilization across all dimensions

Conclusions

Conclusions

This thesis aims to foresee some features that an ideal fully-INP aware Resource Manager should have, as well as some open problems in the INP resource management field.

Fully INP-aware RM features



- Conjunct placement of server and switch resources
- INP alternatives



Open problems


- Accurately determine STC values for all INP solutions
- Determine the number of needed switch tasks for INP solutions
- Differentiate INP solutions based on their life cycle (e.g., short-term batch jobs vs. long-term services)

Questions?

Thank you

-  R. L. Graham, D. Bureddy, P. Lui, H. Rosenstock, G. Shainer, G. Bloch, D. Goldenberg, M. Dubman, S. Kotchubievsky, V. Koushnir, et al.
Scalable hierarchical aggregation protocol (sharp): a hardware architecture for efficient data reduction.
In Proceedings of the First Workshop on Optimization of Communication in HPC, pages 1–10. IEEE Press, 2016.
-  X. Jin, X. Li, H. Zhang, N. Foster, J. Lee, R. Soulé, C. Kim, and I. Stoica.
Netchain: Scale-free sub-rtt coordination.
In 15th USENIX Symposium on Networked Systems Design and Implementation (NSDI 18), pages 35–49, Renton, WA, 2018. USENIX Association.

-  M. Liu, L. Luo, J. Nelson, L. Ceze, A. Krishnamurthy, and K. Atreya.
Incbricks: Toward in-network computation with an in-network cache.
In *Proceedings of the Twenty-Second International Conference on Architectural Support for Programming Languages and Operating Systems, ASPLOS '17*, pages 795–809, New York, NY, USA, 2017. ACM.
-  A. Sapio, I. Abdelaziz, A. Aldilaijan, M. Canini, and P. Kalnis.
In-network computation is a dumb idea whose time has come.
In *Proceedings of the 16th ACM Workshop on Hot Topics in Networks, HotNets-XVI*, pages 150–156, New York, NY, USA, 2017. ACM.

-  M. Schwarzkopf, A. Konwinski, M. Abd-El-Malek, and J. Wilkes.
Omega: flexible, scalable schedulers for large compute clusters.
In *SIGOPS European Conference on Computer Systems (EuroSys)*, pages
351–364, Prague, Czech Republic, 2013.